

LoPA: Scaling dLLM Inference via Lookahead Parallel Decoding

Chenkai Xu^{1*} Yijie Jin^{1*} Jiajun Li² Yi Tu² Guoping Long² Dandan Tu² Mingcong Song² Hongjie Si²
Tianqi Hou² Junchi Yan¹ Zhijie Deng^{1†}

Github: <https://github.com/zhijie-group/LoPA>

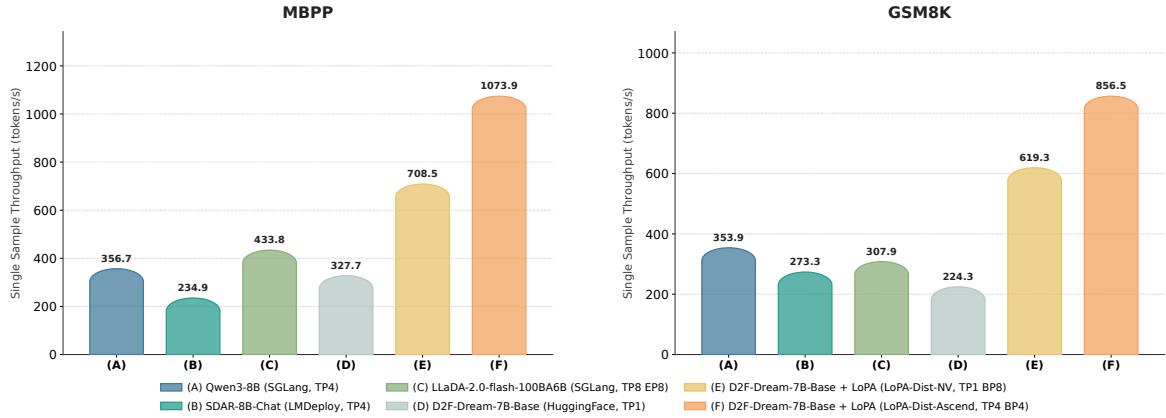


Figure 1. **Throughput performance of LoPA.** LoPA accelerates the **single-sample throughput** for D2F-Dream to up to 1073.9 and 856.5 tokens/s on MBPP and GSM8K respectively, significantly outperforming baselines. More details are provided in Table 7.

Abstract

Diffusion Large Language Models (dLLMs) have demonstrated significant potential for high-speed inference. However, current confidence-driven decoding strategies are constrained by limited parallelism, typically achieving only 1–3 tokens per forward pass (TPF). In this work, we identify that the degree of parallelism during dLLM inference is highly sensitive to the Token Filling Order (TFO). Then, we introduce **Lookahead PArallel Decoding (LoPA)**, a training-free, plug-and-play algorithm, to identify a superior TFO and hence accelerate inference. LoPA concurrently explores distinct candidate TFOs via parallel branches, and selects the one with the highest potential for future parallelism based on branch confidence. We apply LoPA to the state-of-the-art D2F model and observe a substantial enhancement in decoding efficiency. Notably, LoPA increases the TPF of D2F-Dream to 10.1 on the GSM8K while main-

taining performance superior to the Dream baseline. Furthermore, to facilitate this unprecedented degree of parallelism, we develop a specialized multi-device inference system featuring Branch Parallelism (BP), which achieves a single-sample throughput of **1073.9** tokens per second under multi-GPU deployment.

1. Introduction

Diffusion Large Language Models (dLLMs) (Nie et al., 2025; Ye et al., 2025; Gong et al., 2025; Wang et al., 2025b; Cheng et al., 2025) have emerged as a highly promising paradigm for text generation. By iteratively refining a full-mask sequence into text tokens, dLLMs decouple generation depth from sequence length, theoretically offering superior potential for inference speed. Recent studies (Wang et al., 2025b; Cheng et al., 2025) have demonstrated speed surpassing autoregressive (AR) models, thereby providing higher throughput for latency-sensitive applications.

Despite this potential, the practical parallelism of state-of-the-art dLLMs remains constrained. Leading models such as Fast-dLLM (Wu et al., 2025), D2F (Wang et al., 2025b), and SDAR (Cheng et al., 2025) employ confidence-driven sam-

¹Shanghai Jiao Tong University ²Huawei. Correspondence to: Zhijie Deng <zhijied@sjtu.edu.cn>.

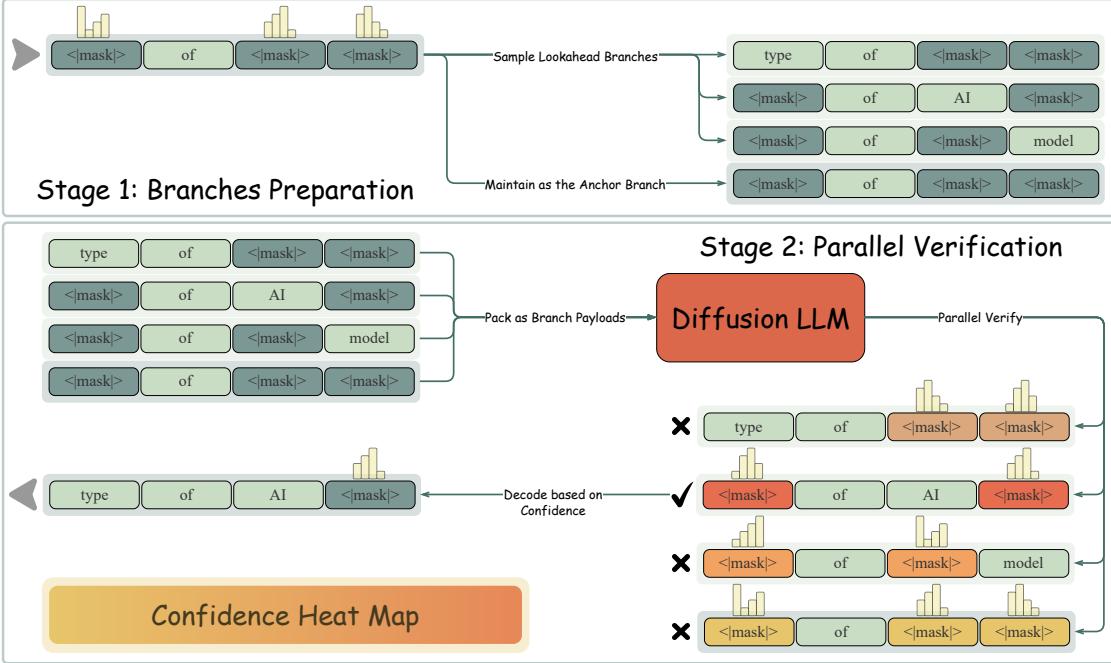


Figure 2. Overview of Lookahead Parallel Decoding (LoPA). In each iteration, LoPA generates a anchor branch alongside multiple lookahead branches (e.g., B_1, \dots, B_k) by independently sampling high-confidence positions from the baseline’s unfilled set. A branch confidence verification mechanism then evaluates all branches in parallel within a single forward pass, selecting the optimal path to maximize future parallelism.

pling, which fills tokens exceeding a confidence threshold τ in each iteration. However, this strategy typically yields only 1–3 tokens per forward pass (TPF) on common tasks such as mathematics and coding, failing to fully unleash the parallel potential of dLLMs.

Our investigation traces this limitation to a fundamental sensitivity: parallelism is bounded by prediction confidence and the confidence is heavily influenced by the Token Filling Order (TFO). As observed in LLaDA (Nie et al., 2025), varying TFOs significantly shifts the generative distribution and confidence landscape. Consequently, the standard strategy of greedily prioritizing positions with the highest current confidence may lead to suboptimal trajectories, raising the question: *Can we actively explore superior TFOs to maximize future confidence and unlock higher parallelism?*

To this end, we introduce **Lookahead PArallel Decoding (LoPA)**, a training-free, plug-and-play algorithm designed to search for TFOs with high future parallelization potential. LoPA operates in three phases per iteration: (1) advancing decoding by sampling an anchor branch B_0 via standard confidence-driven strategies; (2) exploring distinct TFOs beyond B_0 by generating k lookahead branches, each sampling from the top- k high-confidence candidate positions to ensure reliable exploration coverage (Wu et al., 2025); and (3) identifying the optimal path by verifying all $k + 1$

branches in a single forward pass to retain the one with the highest future parallelization potential. By iteratively selecting optimal branches, LoPA can boost overall TPF.

We integrate LoPA with D2F (Wang et al., 2025b), scaling the TPF of D2F-Dream to 10.1 on GSM8K (Cobbe et al., 2021) while maintaining performance scores surpassing the original Dream baseline, and scaling D2F-DiffuCoder to 8.3 on HumanEval+ (Chen et al., 2021; Liu et al., 2023) with marginal performance degradation. To support LoPA, we co-design a multi-device inference system that distributes branches across devices on multiple platforms, achieving a single-sample throughput of **1073.86** tokens/s. We further validate LoPA’s generalizability by integrating it with Vanilla Dream (Ye et al., 2025). Evaluations demonstrate that LoPA effectively scales dLLM parallelism, establishing a clear, controllable speed-accuracy trade-off.

Our contributions are summarized as follows:

- We identify TFO as a key factor influencing dLLM parallelism and propose LoPA, a training-free algorithm that looks ahead to optimize TFO.
- We demonstrate that LoPA scales the TPF of D2F-Dream to 10.1 on GSM8K and D2F-DiffuCoder to 8.3 on HumanEval+, while maintaining comparable performance.

- We develop a specialized Branch Parallel inference system, achieving near-linear scalability and a throughput of **1073.86** tokens/s.

2. Related Work

Diffusion Large Language Models (dLLMs). Autoregressive (AR) models (Achiam et al., 2023; Touvron et al., 2023; Jiang et al., 2023; Liu et al., 2024) have long dominated text generation, yet their sequential decoding imposes an inherent latency bottleneck. To address this, dLLMs (Nie et al., 2025; Ye et al., 2025; Gong et al., 2025; Lou et al., 2023) have emerged as a non-autoregressive paradigm. By iteratively denoising fully masked sequences, dLLMs enable parallel token prediction and leverage bidirectional attention for holistic context modeling. Recent scaling efforts, whether training from scratch (Nie et al., 2025) or initializing from pre-trained AR weights (Ye et al., 2025), have yielded dLLMs with performance competitive to state-of-the-art AR models, validating their potential for high-quality, parallel generation.

Acceleration of dLLMs. Despite their parallel nature, dLLM inference remains computationally expensive due to multi-step denoising and incompatibility with standard KV caching. Acceleration strategies fall into two primary paradigms. Training-based methods compress sampling steps or restructure generation; for instance, dParallel (Chen et al., 2025) applies consistency distillation, while D2F (Wang et al., 2025b) employs asymmetric distillation to enable block-autoregressive pipelining. Training-free methods optimize inference without weight updates. One avenue adapts KV caching to bidirectional attention via approximate schemes (Liu et al., 2025; Wu et al., 2025; Ma et al., 2025). Another focuses on heuristic decoding optimizations, where works like Fast-dLLM (Wu et al., 2025), Prophet (Li et al., 2025), and Credit Decoding (Wang et al., 2025a) exploit confidence patterns or early-layer determinism to skip redundant steps.

Speculative Decoding. Speculative decoding, a standard for AR acceleration (Leviathan et al., 2023; Chen et al., 2023; Spector & Re, 2023), employs efficient draft models for parallel verification. Innovations include tree-based drafting (Miao et al., 2023; Li et al., 2024), multi-head structures (Cai et al., 2024), and draft-free fixed-point iterations (Fu et al., 2024; Zhang et al., 2024). In the dLLM domain, speculative concepts have evolved from relying on external AR guidance (Israel et al., 2025) to self-verification methods like Spiffy (Agrawal et al., 2025) and Free Draft-and-Verification (Wu & Zhang, 2025). While the latter achieve lossless acceleration by maximizing token acceptance within a fixed generation distribution, our approach fundamentally diverges. Instead of passively verifying a static sequence, LoPA actively explores TFO to discover

Algorithm 1 Lookahead Parallel Decoding (LoPA)

```

Input: Sequence  $x_t$ , Mask  $M_t$ , Branch budget  $k$ 
Output: Updated Sequence  $x_{t+1}$ , Mask  $M_{t+1}$ 
// 1. Anchor Branch Construction
    Compute distribution  $p_\theta(\cdot|x_t)$  and confidence scores
    Determine anchor fill set  $I_{fill}$  via Eq. 1 to form branch  $B_0$ 
    Identify remaining unfilled set  $M_{B_0} \leftarrow M_t \setminus I_{fill}$ 
// 2. Lookahead Branches Spawning
    Select top- $k$  positions  $\{p_1, \dots, p_k\} \subset M_{B_0}$  with highest confidence
    for  $j = 1$  to  $k$  do
        Construct branch  $B_j$  by independently sampling position  $p_j$ 
    end for
// 3. Parallel Verification
    Concatenate  $\{B_0, \dots, B_k\}$  as a batch
    Compute scores  $C(B_j)$  via Eq. 2 (Single Pass)
    Select optimal branch  $B^* \leftarrow \arg \max_j C(B_j)$ 
Return Update  $x_{t+1}, M_{t+1}$  according to  $B^*$ 

```

trajectories with superior future confidence, effectively optimizing the output distribution to unlock parallelism beyond standard greedy limits.

3. Methodology

This section first explains the foundational Confidence-Driven Sampling used in regular dLLM inference (Wu et al., 2025) and then elaborates on LoPA.

3.1. Preliminary: Confidence-Driven Sampling for dLLMs

Confidence-driven sampling is a prevalent paradigm for current dLLMs to boost parallelism, which has been widely adopted in advanced methods such as Fast-dLLM (Wu et al., 2025), D2F (Wang et al., 2025b), and SDAR (Cheng et al., 2025). Specifically, given a sequence x_t with a set of masked positions M_t , the dLLM model p_θ outputs a predictive distribution $p_\theta(\cdot|x_t)$. A candidate sequence $\hat{x}_0 \sim p_\theta(\cdot|x_t)$ is sampled, and a confidence function, $\text{Conf}(\cdot)$, assigns a score to each position $i \in M_t$. The set of positions to fill, I_{fill} , is then determined as:

$$S_{\text{high}} = \{i \in M_t \mid \text{Conf}(i) > \tau\}$$

$$I_{fill} = \begin{cases} S_{\text{high}} & \text{if } S_{\text{high}} \neq \emptyset \\ \{\arg \max_{i \in M_t} \text{Conf}(i)\} & \text{otherwise} \end{cases} \quad (1)$$

The algorithm then accepts the predictions according to I_{fill} and moves to the next iteration.

3.2. Lookahead PArallel Decoding (LoPA)

As shown in Figure 2, LoPA addresses the limitation of fixed sampling by looking ahead at multiple TFOs in every

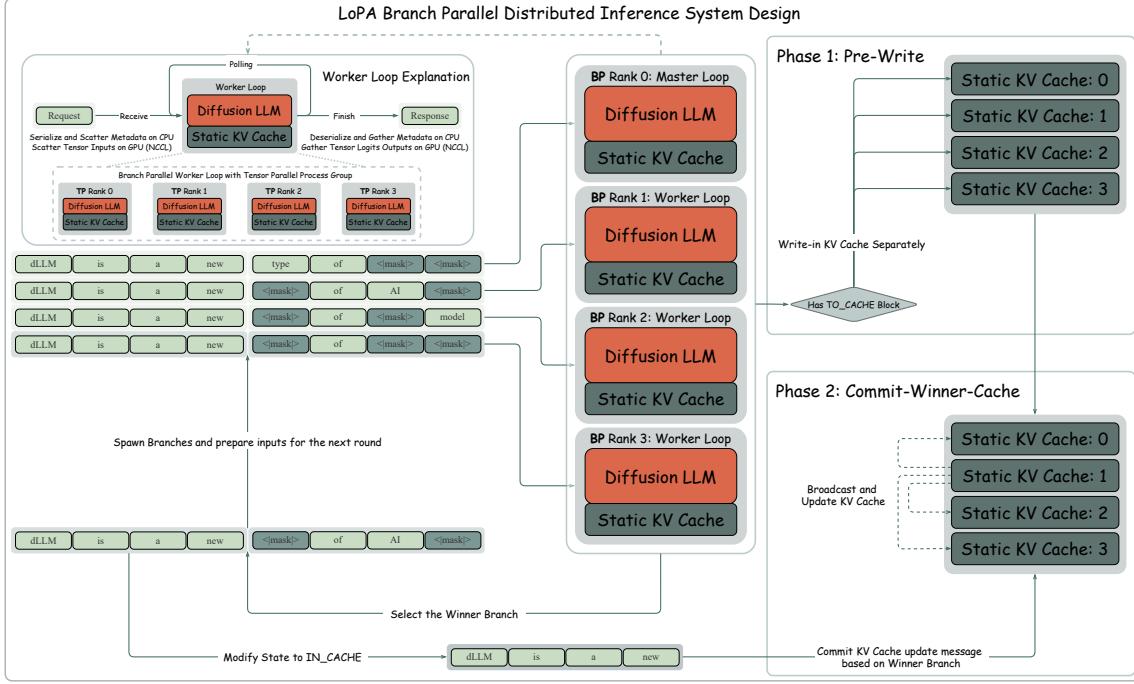


Figure 3. Overview of LoPA Branch Parallel Distributed Inference System Design. A key distinction lies in the KV cache management protocol tailored for different backends: **LoPA-Dist-NV** utilizes a robust two-phase update mechanism to ensure consistency, whereas **LoPA-Dist-Ascend** adopts a streamlined single-phase update strategy for optimized serving efficiency.

decoding iteration. It generates multiple sampling branches concurrently and identifies the one with superior potential for parallel decoding. The detailed procedure is outlined in Algorithm 1.

Look ahead Multiple TFOs in Parallel. Standard confidence-driven sampling relies on a single anchor branch where only positions in I_{fill} are sampled. LoPA extends this by exploring one step further. To ensure exploration is both effective and reliable, we prioritize sampling tokens with higher confidence, a strategy that has been proved in Fast-dLLM (Wu et al., 2025) to yield more stable predictions. Specifically, in addition to the anchor branch B_0 , we generate k competitive branches. We identify the top- k positions from the unfilled set M_{B_0} that possess the highest confidence scores. For each identified position, we sample it independently to create a distinct branch. This results in a set of k new branches $\{B_1, \dots, B_k\}$, each possessing a unique partially filled sequence x_{B_j} and a unfilled set M_{B_j} .

Branch Confidence-based Verification. Inspired by DeepConf (Fu et al., 2025), we design a branch confidence metric to guide the selection among candidate decoding paths. Formally, the confidence of a branch B_j is defined as the average prediction confidence over its remaining unfilled

positions M_{B_j} :

$$C(B_j) = \frac{1}{|M_{B_j}|} \sum_{i \in M_{B_j}} \text{Conf}(i) \quad (2)$$

A higher branch confidence indicates that more unfilled positions are likely to be accepted in the next decoding step. This directly increases the number of tokens filled per iteration, thereby enhancing the overall parallelism. Beyond this mean confidence, branch confidence can also be quantified by other methods (Fu et al., 2025), such as applying a sliding window to assess local quality or averaging confidence over the least confident segment to identify weak links. We adopt this naive average confidence as the default for its simplicity and robust performance.

This branch confidence verification mechanism offers distinct advantages. First, it facilitates the packing and verification of all candidate branches within a single forward pass. Second, the logits computed during evaluation are repurposed for the subsequent decoding step, obviating the need for additional computation.

3.3. Application: Integration with D2F

LoPA integrates seamlessly with D2F (Wang et al., 2025b) by treating all active blocks as a single window for branch exploration. Crucially, within this window, we replace the original block-level causal attention with full attention.

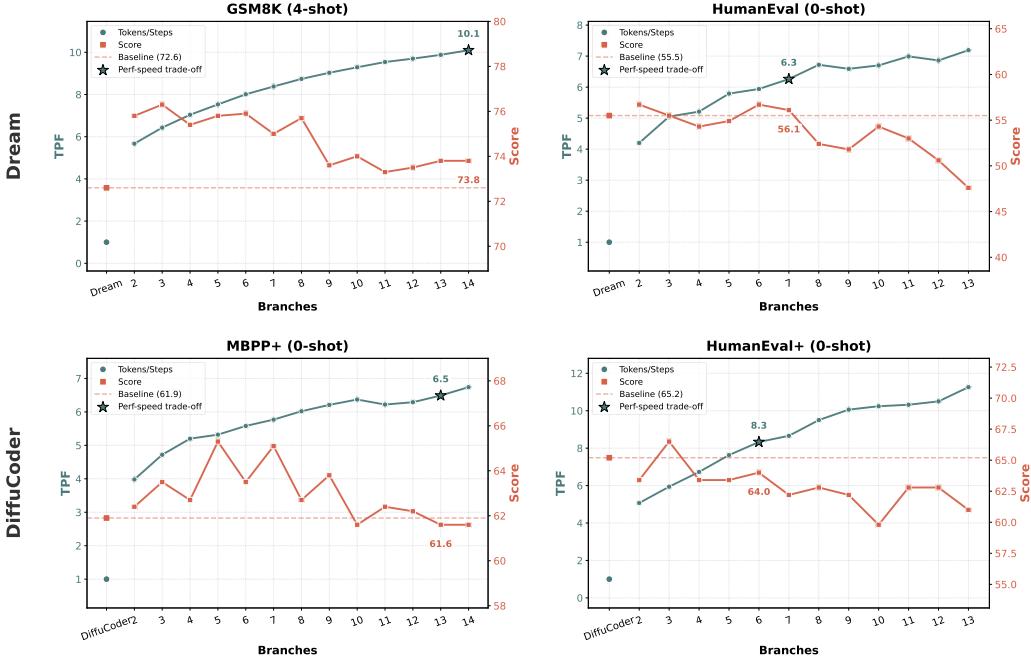


Figure 4. Scaling Curves of LoPA. LoPA scales the tokens per forward pass (TPS) for D2F-Dream and D2F-DiffuCoder to up to 10.1 and 8.3 on GSM8k and HumanEval+ respectively, with comparable performance.

This design yields two key benefits. First, the simplified attention mechanism significantly reduces complexity and enhances compatibility with mainstream inference frameworks, directly boosting computational speed. Second, we empirically find that this localized full attention maintains or even improves generation quality. By enabling blocks to attend to “future” tokens within the limited window, the information flow is enriched without disrupting the global causal dependency required for valid generation.

3.4. System Implementation

To fully unleash the parallelism inherent in LoPA, we introduce **LoPA-Dist**, a high-throughput distributed inference system co-designed with the LoPA algorithm pipeline. LoPA-Dist introduces *Branch Parallelism (BP)* to distribute candidate branches across multiple computing devices, orchestrating synchronized execution to maximize hardware utilization. We provide two specialized implementations of LoPA-Dist tailored for different hardware ecosystems and deployment scenarios.

LoPA-Dist-NV: Latency-Oriented Optimization on CUDA. Targeting the NVIDIA CUDA platform, we developed **LoPA-Dist-NV**, a specialized implementation optimized for ultra-low latency single-sample acceleration. This system employs a pre-allocated static KV cache tightly coupled with the model architecture, effectively circumventing the overhead of dynamic object instantiation and the costly scattering of KV cache tensors via NCCL. To maintain

context consistency across divergent branches without sacrificing speed, LoPA-Dist-NV implements a novel two-phase update protocol.

As illustrated in Figure 3, during the forward pass, the system executes a *Pre-Write* phase: each device speculatively writes the features of its TO_CACHE blocks directly into designated cache slots. Since branch divergence leads to temporary cache inconsistency, we execute a *Commit-Winner-Cache* phase: once the optimal branch is identified, its corresponding KV cache features are broadcast to all peer devices, overwriting local entries to enforce global synchronization. To further minimize runtime latency, LoPA-Dist-NV integrates FlashAttention (Dao et al., 2022) backends via SDPA, pre-merges LoRA weights, utilizes fused kernels for RMS Norm and RoPE, and implements an approximate prefix caching mechanism.

LoPA-Dist-Ascend: High-Performance Serving on Ascend 910C. Building upon the foundation of LoPA-Dist-NV, we architected **LoPA-Dist-Ascend**, a high-throughput serving engine specifically optimized to exploit the computational power of the Huawei Ascend 910C. While retaining the core algorithmic logic of the CUDA implementation, LoPA-Dist-Ascend adopts a vLLM-like (Kwon et al., 2023) architecture with a hybrid parallelism strategy orchestrated via process groups. Specifically, we assign four NPUs to each branch (TP4) to accelerate single-pass forward propagation through Tensor Parallelism, while scaling throughput via Branch Parallelism across groups.

Model	Decoding Algo	MBPP 3-shot		Math 4-shot		HumanEval 0-shot		GSM8K 4-shot	
		TPF	Score	TPF	Score	TPF	Score	TPF	Score
Dream	Vanilla	1.0	56.2	1.0	33.7	1.0	55.5	1.0	72.6
	Fast-dLLM	1.9	55.6	1.9	37.6	1.8	55.5	2.1	72.6
	LoPA	3.3	54.8	3.4	37.0	2.9	53.0	3.1	73.3
D2F-Dream	Vanilla	2.3	53.8	2.6	36.8	2.5	56.1	3.1	78.5
	LoPA	5.4	56.0	8.0	35.2	6.3	56.1	10.1	73.8

Table 1. Accuracy-preserving parallelism scaling of Dream on multiple benchmarks across multiple branches. TPF denotes Tokens Per Forward pass. LoPA significantly scales the TPF of D2F-Dream while maintaining or exceeding baseline scores.

Model	Decoding Algo	MBPP+		HumanEval+	
		TPF	Score	TPF	Score
DiffuCoder	Vanilla	1.0	61.9	1.0	65.2
D2F-DiffuCoder	Vanilla	2.2	61.9	2.2	65.9
	LoPA	6.7	61.6	8.3	64.0

Table 2. Accuracy-preserving parallelism scaling of DiffuCoder on MBPP+ and HumanEval+ benchmarks. LoPA boosts TPF by nearly 4× compared to the vanilla D2F baseline with minimal impact on generation quality.

To streamline system complexity for large-scale serving, LoPA-Dist-Ascend utilizes a block-wise causal mask for the attention mechanism. This design choice eliminates the necessity for the asynchronous *Commit-Winner-Cache* phase; instead, it relies solely on a direct *Pre-Write* operation, which inherently maintains inter-branch consistency through masking. To achieve optimal throughput, LoPA-Dist-Ascend integrates a suite of hardware-aware optimizations:

- *NPU-Aware FlashAttention & Pipelining*: Addressing the self-attention bottleneck, we implemented a custom FlashAttention (Dao et al., 2022) kernel optimized for the Ascend 910C’s Cube Units. We further integrate advanced pipelining techniques (Song et al., 2025) to effectively mask latencies between Cube and Vector operations. By tiling Q, K, and V matrices into the on-chip SRAM, we minimize HBM data movement, achieving near-theoretical IO throughput.
- *Graph Compilation & Operator Fusion*: Since dLLMs involve iterative denoising steps necessitating repetitive kernel launches, we utilize graph compilation to fuse element-wise operations (e.g., bias addition and activations), significantly reducing launch overhead.
- *Diffusion-Tailored Memory Management*: We implement a paged memory management system similar to PagedAttention but specifically tailored for the state-space requirements of diffusion steps.

- *System-Level Optimizations*: The engine further incorporates W8A8 quantization, QKV merging, and fully asynchronous inference scheduling.

4. Experiments

4.1. Experimental Settings

Our experiments primarily focus on the D2F model (Wang et al., 2025b), which is the first open-source dLLM whose inference throughput surpasses that of autoregressive (AR) models. Specifically, we integrate LoPA with the Dream-Instruct-7B (Ye et al., 2025) and DiffuCoder-Instruct-7B (Gong et al., 2025) models trained with the D2F objective, denoted as D2F-Dream and D2F-DiffuCoder, respectively. To further verify the generalizability of our method across different architectures, we also integrate LoPA into the vanilla Dream-Instruct-7B model utilizing confidence-based decoding.

Regarding the hardware configuration, we adopt distinct setups for theoretical analysis and system-level evaluation. For the theoretical speedup experiments, we target a distributed setting with 8 NVIDIA A100 GPUs utilizing standard task parallelism. Note that the performance metrics in this phase are projected based on single-GPU benchmarks to derive theoretical upper bounds. For the system-level acceleration experiments, we deploy our method on two distinct high-performance platforms: (1) for the CUDA platform, we utilize 8 NVIDIA H200 GPUs to evaluate performance under Branch Parallelism (BP) configurations of 4 and 8

Model	Platform	MBPP				GSM8K			
		Avg TPS	Max TPS	TPF	Latency	Avg TPS	Max TPS	TPF	Latency
D2F-Dream-Base	LoPA-Dist-NV	708.48	1470.95	15.55	0.74	619.33	1299.25	13.16	0.85
	LoPA-Dist-Ascend	1073.86	2400.12	11.92	0.78	856.46	2751.61	9.34	0.75
D2F-Dream-Instruct	LoPA-Dist-NV	636.55	1811.71	9.52	0.14	609.90	1407.56	11.42	0.26
	LoPA-Dist-Ascend	896.21	2586.73	8.64	0.11	897.10	1868.16	9.30	0.21

Table 3. System performance of LoPA. The results demonstrate that our system efficiently translates algorithmic parallelism (high TPF) into significant wall-clock acceleration, achieving average throughputs exceeding 1000 tokens/s on the specialized LoPA-Dist-Ascend engine.

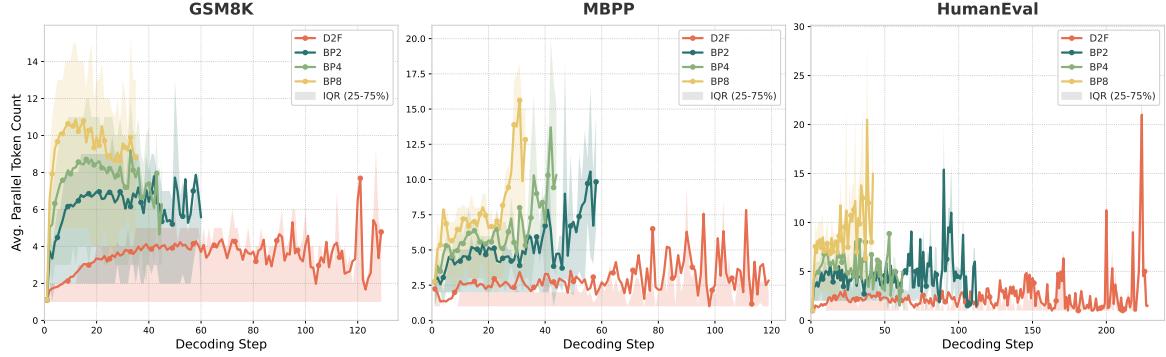


Figure 5. Scaling analysis of LoPA on D2F-Dream with varying branch counts. The results illustrate that LoPA effectively scales the TPF of D2F to a peak exceeding 10, thereby significantly reducing the total number of decoding steps.

(BP4 and BP8); (2) for the Ascend platform, we employ a cluster of 8 Ascend 910C NPUs, configured to support a hybrid parallelism strategy combining Tensor Parallelism (TP4) and Branch Parallelism (BP4).

4.2. Main Results

Baselines. We compare LoPA against the vanilla dLLM baseline. Additionally, we include Fast-dLLM (Wu et al., 2025) as a competitive accelerated baseline and the standard D2F model (Wang et al., 2025b) to isolate the performance gains attributable to our lookahead strategy.

Benchmarks. Our evaluation spans diverse domains including mathematical reasoning and code generation. For mathematical reasoning, we employ GSM8K (Cobbe et al., 2021) and MATH (Hendrycks et al., 2021). For code generation, we utilize HumanEval (Chen et al., 2021) and MBPP (Austin et al., 2021). Furthermore, to ensure a rigorous and robust assessment, we include the extended benchmarks HumanEval+ and MBPP+ from EvalPlus (Liu et al., 2023).

Scaling D2F Inference. We analyze the speed-quality trade-off by varying branch count k . As detailed in Figure 4, while increasing k enhances parallelism, excessive branching may induce fluctuations by prioritizing future confidence. Consequently, with an optimal k , LoPA substantially scales the inference speed. Specifically, it increases the TPF of D2F-Dream to 10.1 on GSM8K while maintaining performance

superior to the Dream baseline, and scales D2F-DiffuCoder to 8.3 on HumanEval+ with marginal performance degradation. Tables 1 and 2 confirm this efficacy. For instance, on MATH, LoPA attains a high TPF of 8.0 while achieving a performance score superior to the Vanilla baseline. This validates that LoPA enables high-parallelism decoding without compromising generation quality.

Generalizability Verification. Beyond D2F-based architectures, LoPA demonstrates strong universality. As shown in Table 1, when applied to the Vanilla Dream utilizing confidence-based decoding, LoPA scales the TPF to 3.4 on Math while maintaining comparable performance. This evidence validates LoPA as a generalized, plug-and-play scaling solution applicable to broad confidence-driven dLLMs.

System Throughput. To fully unleash the parallelism enabled by LoPA across diverse hardware backends, we evaluate the wall-clock performance using our co-designed Branch Parallel (BP) inference system, as illustrated in Figure 3. Under multi-device deployment, the system exhibits near-linear scalability with respect to the branch count, effectively translating high TPF into tangible throughput gains. As visualized in Figure 1, the system achieves a single-sample throughput of **1073.86** tokens per second, where both the D2F baseline and D2F + LoPA employ identical, slightly lowered decoding thresholds to prioritize and maximize inference speed.

LoPA: Scaling dLLM Inference via Lookahead Parallel Decoding

Model	Sys. Arch.	Settings	MBPP 3-shot				GSM8K 4-shot			
			Avg TPS	Max TPS	Top-10 TPS	Score	Avg TPS	Max TPS	Top-10 TPS	Score
LoPA-Dist-NV	D2F-Dream-Base	S1	415.19	813.04	720.35	53.00	345.52	959.05	704.39	75.97
		S2	500.33	1185.77	874.87	53.40	402.52	913.12	842.83	73.54
		S3	550.37	1472.41	929.72	51.20	436.22	994.82	885.27	71.19
		S4	589.22	1576.93	1006.57	47.20	475.58	1203.61	1028.15	68.16
		S5	633.16	1408.40	963.67	46.80	516.85	1212.65	1055.08	66.79
		S6	678.26	1615.30	1150.65	41.80	546.72	1225.21	1121.57	64.14
LoPA-Dist-Ascend	D2F-Dream-Base	S7	466.27	784.33	764.52	51.80	416.91	909.82	841.95	71.27
		S8	545.90	1497.22	927.67	51.40	486.94	1176.14	959.37	68.39
		S9	588.00	1584.28	983.09	48.60	520.70	1250.67	1056.01	68.01
		S10	637.38	1552.56	1028.97	47.00	558.01	1115.26	1071.66	65.05
		S11	655.45	1535.10	1059.72	43.80	592.94	1315.93	1155.11	64.44
		S12	708.48	1470.95	1132.78	39.80	619.33	1299.25	1201.18	60.88
LoPA-Dist-NV	D2F-Dream-Instruct	S13	615.74	2173.7	1253.07	50.20	492.94	1337.60	1158.18	75.06
		S14	753.78	2115.55	1397.85	50.20	589.77	1532.99	1342.79	72.86
		S15	842.97	2470.79	1538.16	50.00	644.34	1723.19	1476.24	70.58
		S16	923.35	2647.12	1513.54	45.60	700.14	1756.58	1601.93	68.69
		S17	994.88	2740.54	1739.85	43.00	754.75	2583.76	1848.82	64.29
		S18	1073.86	2400.12	1939.22	41.80	856.46	2751.61	2098.72	62.55
LoPA-Dist-Ascend	D2F-Dream-Instruct	S1	305.74	959.00	695.88	52.80	330.62	758.34	674.53	78.17
		S2	373.23	1302.99	877.12	51.40	402.63	961.29	804.31	74.22
		S3	451.62	1419.09	1143.30	53.00	444.73	943.22	870.85	73.39
		S4	503.71	1779.60	1226.72	46.60	495.93	1131.64	941.23	72.48
		S5	568.65	1660.89	1317.38	42.00	540.76	1185.14	1033.60	68.99
		S6	615.95	1951.86	1542.82	37.60	568.75	1352.22	1139.06	65.88
LoPA-Dist-NV	D2F-Dream-Instruct	S7	325.15	697.49	620.42	50.80	379.42	839.65	710.10	75.28
		S8	408.37	1182.69	866.90	51.00	449.56	934.55	838.35	75.13
		S9	465.55	1097.40	1016.91	50.60	497.47	1172.31	946.98	74.75
		S10	544.72	1542.99	1145.55	46.80	539.28	1147.95	1021.96	71.34
		S11	591.57	1578.00	1204.05	42.20	580.04	1292.18	1132.19	66.94
		S12	636.55	1811.71	1500.59	36.00	609.90	1407.56	1159.28	65.50
LoPA-Dist-Ascend	D2F-Dream-Instruct	S13	412.90	911.73	911.73	50.80	515.01	1235.84	1090.45	76.12
		S14	525.66	1546.34	1143.37	48.40	619.58	1424.32	1310.35	75.36
		S15	625.53	1729.78	1435.06	46.20	689.89	1644.74	1356.36	72.63
		S16	716.19	1780.41	1558.00	43.80	770.78	1589.69	1480.56	71.49
		S17	796.65	1798.14	1687.69	39.80	837.21	1782.80	1517.90	67.78
		S18	896.21	2586.73	2086.04	36.40	897.10	1868.16	1642.72	66.87

Table 4. Performance ablation study of D2F-Dream models on different platforms, corresponding to settings S1-S18. The results illustrate the trade-off between inference throughput and generation quality across varying branch configurations and system backends.

The comprehensive evaluation results are summarized in Table 3, where the optimal configurations selected for peak performance are TP1+BP8 for LoPA-Dist-NV and TP4+BP4 for LoPA-Dist-Ascend. It is important to note that this table presents the peak system throughput achieved under the optimal hyperparameter configurations mentioned above. For a granular analysis of performance across different settings (e.g., varying decoding thresholds and parallel strategies), please refer to the detailed ablation study provided in Table 4.

It is noteworthy that we provide a cross-platform implementation to demonstrate the versatility of our system. While both implementations adhere to the unified architectural design shown in Figure 3, they employ distinct backend strategies: the **CUDA** results correspond to a naive distributed implementation, whereas the **Ascend** results are derived from our specialized inference engine, which incorpo-

rates PagedAttention-like (Kwon et al., 2023) mechanisms adapted for deployment on Ascend devices. In addition to the standard scores, we detail the system performance metrics including the Average TPS, which reflects the sustained generation speed; the Max TPS, demonstrating the peak throughput capability; TPF, which quantifies the efficiency of parallel decoding; and Latency, which provides a direct measure of wall-clock inference speed.

Scaling Analysis. As shown in Figure 5, we conduct a comprehensive scaling analysis of LoPA on D2F-Dream with varying branch counts across multiple benchmarks. The results illustrate that LoPA effectively scales the TPF of D2F to a peak exceeding 10, thereby significantly reducing the total number of decoding steps. We further observe that math tasks like GSM8K exhibit high parallelism in the middle stages of generation, whereas code tasks such as MBPP and HumanEval show higher parallelism in the later

stages.

5. Conclusion

In this paper, we propose **Lookahead Parallel Decoding (LoPA)**, a training-free algorithm designed to break the parallelism bottleneck in dLLM inference by identifying the optimal Token Filling Order (TFO). By concurrently exploring candidate branches to maximize future confidence, LoPA significantly scales decoding efficiency, boosting the Tokens Per Forward pass (TPF) of D2F-Dream to **10.1** on GSM8K and D2F-DiffuCoder to **8.3** on HumanEval+ while preserving generation quality. Furthermore, we facilitate these gains with a specialized *Branch Parallel (BP)* inference system that ensures substantial wall-clock speedups, establishing LoPA as a robust solution for efficient non-autoregressive sequence generation.

References

- Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., Anadkat, S., et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Agrawal, S., Garrepalli, R., Goel, R., Lee, M., Lott, C., and Porikli, F. Spify: Multiplying diffusion llm acceleration via lossless speculative decoding. *arXiv preprint arXiv:2509.18085*, 2025.
- Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Cai, T., Li, Y., Geng, Z., Peng, H., Lee, J. D., Chen, D., and Dao, T. Medusa: Simple llm inference acceleration framework with multiple decoding heads. *arXiv preprint arXiv:2401.10774*, 2024.
- Chen, C., Borgeaud, S., Irving, G., Lespiau, J.-B., Sifre, L., and Jumper, J. Accelerating large language model decoding with speculative sampling. *arXiv preprint arXiv:2302.01318*, 2023.
- Chen, M., Tworek, J., Jun, H., Yuan, Q., de Oliveira Pinto, H. P., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., Ray, A., Puri, R., Krueger, G., Petrov, M., Khlaaf, H., Sastry, G., Mishkin, P., Chan, B., Gray, S., Ryder, N., Pavlov, M., Power, A., Kaiser, L., Bavarian, M., Winter, C., Tillet, P., Such, F. P., Cummings, D., Plappert, M., Chantzis, F., Barnes, E., Herbert-Voss, A., Guss, W. H., Nichol, A., Paino, A., Tezak, N., Tang, J., Babuschkin, I., Balaji, S., Jain, S., Saunders, W., Hesse, C., Carr, A. N., Leike, J., Achiam, J., Misra, V., Morikawa, E., Radford, A., Knight, M., Brundage, M., Murati, M., Mayer, K., Welinder, P., McGrew, B., Amodei, D., McCandlish, S., Sutskever, I., and Zaremba, W. Evaluating large language models trained on code. 2021.
- Chen, Z., Fang, G., Ma, X., Yu, R., and Wang, X. dparallel: Learnable parallel decoding for dllms. *arXiv preprint arXiv:2509.26488*, 2025.
- Cheng, S., Bian, Y., Liu, D., Zhang, L., Yao, Q., Tian, Z., Wang, W., Guo, Q., Chen, K., Qi, B., and Zhou, B. Sdar: A synergistic diffusion-autoregression paradigm for scalable sequence generation, 2025. URL <https://arxiv.org/abs/2510.06303>.
- Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., Hesse, C., and Schulman, J. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Dao, T., Fu, D. Y., Ermon, S., Rudra, A., and Ré, C. Flashattention: Fast and memory-efficient exact attention with io-awareness, 2022. URL <https://arxiv.org/abs/2205.14135>.
- Fu, Y., Bailis, P., Stoica, I., and Zhang, H. Break the sequential dependency of llm inference using lookahead decoding. *arXiv preprint arXiv:2402.02057*, 2024.
- Fu, Y., Wang, X., Tian, Y., and Zhao, J. Deep think with confidence, 2025. URL <https://arxiv.org/abs/2508.15260>.
- Gong, S., Zhang, R., Zheng, H., Gu, J., Jaitly, N., Kong, L., and Zhang, Y. Diffucoder: Understanding and improving masked diffusion models for code generation, 2025. URL <https://arxiv.org/abs/2506.20639>.
- Hendrycks, D., Burns, C., Kadavath, S., Arora, A., Basart, S., Tang, E., Song, D., and Steinhardt, J. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.
- Israel, D., Broeck, G. V. d., and Grover, A. Accelerating diffusion llms via adaptive parallel decoding. *arXiv preprint arXiv:2506.00413*, 2025.
- Jiang, A. Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D. S., de las Casas, D., Bressand, F., Lengyel, G., Lample, G., Saulnier, L., Lavaud, L. R., Lachaux, M.-A., Stock, P., Scao, T. L., Lavril, T., Wang, T., Lacroix, T., and Sayed, W. E. Mistral 7b, 2023. URL <https://arxiv.org/abs/2310.06825>.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model

- serving with pagedattention, 2023. URL <https://arxiv.org/abs/2309.06180>.
- Leviathan, Y., Kalman, M., and Matias, Y. Fast inference from transformers via speculative decoding. In *International Conference on Machine Learning*, pp. 19274–19286. PMLR, 2023.
- Li, P., Zhou, Y., Muhtar, D., Yin, L., Yan, S., Shen, L., Liang, Y., Vosoughi, S., and Liu, S. Diffusion language models know the answer before decoding. *arXiv preprint arXiv:2508.19982*, 2025.
- Li, Y., Wei, F., Zhang, C., and Zhang, H. Eagle: Speculative sampling requires rethinking feature uncertainty. *arXiv preprint arXiv:2401.15077*, 2024.
- Liu, A., Feng, B., Xue, B., Wang, B., Wu, B., Lu, C., Zhao, C., Deng, C., Zhang, C., Ruan, C., et al. Deepseek-v3 technical report. *arXiv preprint arXiv:2412.19437*, 2024.
- Liu, J., Xia, C. S., Wang, Y., and Zhang, L. Is your code generated by chatgpt really correct? rigorous evaluation of large language models for code generation. *Advances in Neural Information Processing Systems*, 36:21558–21572, 2023.
- Liu, Z., Yang, Y., Zhang, Y., Chen, J., Zou, C., Wei, Q., Wang, S., and Zhang, L. dllm-cache: Accelerating diffusion large language models with adaptive caching. *arXiv preprint arXiv:2506.06295*, 2025.
- Lou, A., Meng, C., and Ermon, S. Discrete diffusion modeling by estimating the ratios of the data distribution. *arXiv preprint arXiv:2310.16834*, 2023.
- Ma, X., Yu, R., Fang, G., and Wang, X. dkv-cache: The cache for diffusion language models. *arXiv preprint arXiv:2505.15781*, 2025.
- Miao, X., Oliaro, G., Zhang, Z., Cheng, X., Wang, Z., Wong, R. Y. Y., Chen, Z., Arfeen, D., Abhyankar, R., and Jia, Z. Specinfer: Accelerating generative llm serving with speculative inference and token tree verification. *arXiv preprint arXiv:2305.09781*, 1(2):4, 2023.
- Nie, S., Zhu, F., You, Z., Zhang, X., Ou, J., Hu, J., Zhou, J., Lin, Y., Wen, J.-R., and Li, C. Large language diffusion models, 2025. URL <https://arxiv.org/abs/2502.09992>.
- Song, M., Tang, X., Hou, F., Li, J., Wei, W., Ma, Y., Xiao, R., Si, H., Jiang, D., Yin, S., Hu, Y., and Long, G. Xy-serve: End-to-end versatile production serving for dynamic llm workloads. In *Proceedings of the 31st ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 1*, ASPLOS ’26, pp. 314–329, New York, NY, USA, 2025. Association for Computing Machinery. ISBN 9798400721656. doi: 10.1145/3760250.3762228. URL <https://doi.org/10.1145/3760250.3762228>.
- Spector, B. and Re, C. Accelerating llm inference with staged speculative decoding. *arXiv preprint arXiv:2308.04623*, 2023.
- Touvron, H., Martin, L., Stone, K., Albert, P., Almahairi, A., Babaei, Y., Bashlykov, N., Batra, S., Bhargava, P., Bhosale, S., et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- Wang, K., Jiang, Z., Feng, H., Zhao, W., Liu, L., Li, J., Lan, Z., and Lin, W. Creditdecoding: Accelerating parallel decoding in diffusion large language models with trace credits. *arXiv preprint arXiv:2510.06133*, 2025a.
- Wang, X., Xu, C., Jin, Y., Jin, J., Zhang, H., and Deng, Z. Diffusion llms can do faster-than-ar inference via discrete diffusion forcing, 2025b. URL <https://arxiv.org/abs/2508.09192>.
- Wu, C., Zhang, H., Xue, S., Liu, Z., Diao, S., Zhu, L., Luo, P., Han, S., and Xie, E. Fast-dllm: Training-free acceleration of diffusion llm by enabling kv cache and parallel decoding, 2025. URL <https://arxiv.org/abs/2505.22618>.
- Wu, S. and Zhang, J. Free draft-and-verification: Toward lossless parallel decoding for diffusion large language models. *arXiv preprint arXiv:2510.00294*, 2025.
- Ye, J., Xie, Z., Zheng, L., Gao, J., Wu, Z., Jiang, X., Li, Z., and Kong, L. Dream 7b: Diffusion large language models, 2025. URL <https://arxiv.org/abs/2508.15487>.
- Zhang, J., Wang, J., Li, H., Shou, L., Chen, K., Chen, G., and Mehrotra, S. Draft& verify: Lossless large language model acceleration via self-speculative decoding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 11263–11282, 2024.

A. Hardware Environment: Huawei Ascend 910C

Diffusion models typically involve iterative denoising steps, leading to repetitive kernel launches. We utilize the graph compilation capabilities of the CANN software stack to fuse element-wise operations (such as bias-add and activation functions) and optimize the computation graph. This reduces the kernel launch overhead significantly during the multi-step diffusion sampling process.

The Ascend inference engine is deployed on the Huawei Ascend 910C AI processor, a next-generation Neural Processing Unit (NPU) designed for large-scale AI training and inference. The Ascend 910C is built upon the advanced Da Vinci architecture, featuring high-performance Cube Units for matrix operations and Vector Units for general-purpose calculations.

Compared to its predecessors, the 910C offers significant improvements in compute density and inter-chip interconnect bandwidth, making it highly suitable for the memory-intensive nature of Diffusion Language Models.

The platform utilizes the Computational Architecture for Neural Networks (CANN), which acts as the bridge between the upper-level deep learning frameworks and the underlying hardware resources, enabling efficient operator mapping and memory management.

B. Detailed Hyperparameter Configurations

In this section, we present the specific hyperparameter settings used for the D2F-Dream LoPA and D2F-DiffuCoder LoPA experiments. Table 5 details the prompt shot counts, D2F parameters (including Block size, τ_{add} , τ_{act} , and τ_{conf}), and the number of LoPA branches for each dataset. These configurations correspond to the evaluation results reported in Table 1 and Table 2.

Dataset	Shots	Block Size	τ_{add}	τ_{act}	τ_{conf}	LoPA Branches
D2F-Dream LoPA						
GSM8K	4-Shot	32	0.1	0.95	0.90	14
MATH	4-Shot	16	0.1	0.95	0.90	7
HumanEval	0-Shot	32	0.3	0.95	0.95	7
MBPP	3-Shot	32	0.3	0.95	0.95	7
D2F-DiffuCoder LoPA						
HumanEval+	0-Shot	32	0.3	0.95	0.95	6
MBPP+	0-Shot	32	0.3	0.95	0.90	14

Table 5. Hyperparameter settings for D2F-Dream LoPA and D2F-DiffuCoder LoPA. The listed parameters correspond to the main results in Table 1 and Table 2.

C. Detailed Baseline Results

To provide a comprehensive performance assessment, we compare our method against representative state-of-the-art models from both the Diffusion LLM and Autoregressive (AR) LLM families. Specifically, we include SDAR (Cheng et al., 2025) as a strong dLLM baseline known for its high generation quality, and Qwen2.5-7B-Instruct as a standard benchmark for AR inference speed and accuracy. Table 7 details these comparisons, highlighting the trade-offs between throughput (TPS), parallelism (TPF), and generation scores.

Settings	Sys. Arch.	Precision	TP Size	BP Size	Block Size	Seq. Length	τ_{add}	τ_{act}	τ_{conf}
S1	LoPA-Dist-NV	BF16	1	4	32	512	0.1	0.95	0.95
S2		BF16	1	4	32	512	0.1	0.9	0.9
S3		BF16	1	4	32	512	0.1	0.85	0.85
S4		BF16	1	4	32	512	0.1	0.8	0.8
S5		BF16	1	4	32	512	0.1	0.75	0.75
S6		BF16	1	4	32	512	0.1	0.7	0.7
S7		BF16	1	8	32	512	0.1	0.95	0.95
S8		BF16	1	8	32	512	0.1	0.9	0.9
S9		BF16	1	8	32	512	0.1	0.85	0.85
S10		BF16	1	8	32	512	0.1	0.8	0.8
S11		BF16	1	8	32	512	0.1	0.75	0.75
S12		BF16	1	8	32	512	0.1	0.7	0.7
S13	LoPA-Dist-Ascend	W8A8	4	4	32	512	0.1	0.95	0.95
S14		W8A8	4	4	32	512	0.1	0.9	0.9
S15		W8A8	4	4	32	512	0.1	0.85	0.85
S16		W8A8	4	4	32	512	0.1	0.8	0.8
S17		W8A8	4	4	32	512	0.1	0.75	0.75
S18		W8A8	4	4	32	512	0.1	0.7	0.7

Table 6. Hyperparameter configurations for each setting employed in the performance ablation study. TP: Tensor Parallel, BP: Branch Parallel.

Model	MBPP				GSM8K			
	Avg TPS	TPF	Latency	Score	Avg TPS	TPF	Latency	Score
D2F-Dream-7B-Base (CUDA)	327.69	5.64	1.91	45.00	224.33	3.90	2.68	64.90
D2F-Dream-7B-Instruct (CUDA)	206.37	3.13	0.54	45.00	247.90	4.01	0.89	69.07
Qwen3-8B (SGLang)	317.41	1.00	0.84	78.92	317.31	1.00	0.49	93.63
SDAR-4B-Chat	–	1.50	–	52.0	–	2.00	–	88.90
SDAR-8B-Chat (LMDeploy)	234.90	1.60	0.25	72.00	273.31	2.10	0.72	91.30
LLaDA2.0-flash (SGLang)	433.81	2.70	0.21	88.30	307.92	5.30	0.64	96.06

Table 7. Performance comparison of baseline results on multiple benchmarks across multiple platforms. In this configuration, D2F hyperparameters are specifically aligned with the peak throughput capabilities of the current LoPA system, where decoding thresholds are adjusted to ensure maximum inference speed. We incorporate SDAR-4B-Chat and Qwen3-8B to benchmark against SOTA dLLMs and AR models. TPS denotes Tokens Per Second. The Ascend engine demonstrates the effectiveness of our system design compared to the naive CUDA implementation.